

CS 6023 - GPU Programming

Background: CPU Architecture

18/01/2019

Agenda

How do we contrast CPUs from GPUs?


Part 1 of 2:

Today: CPU architecture background

Part 2 of 2:

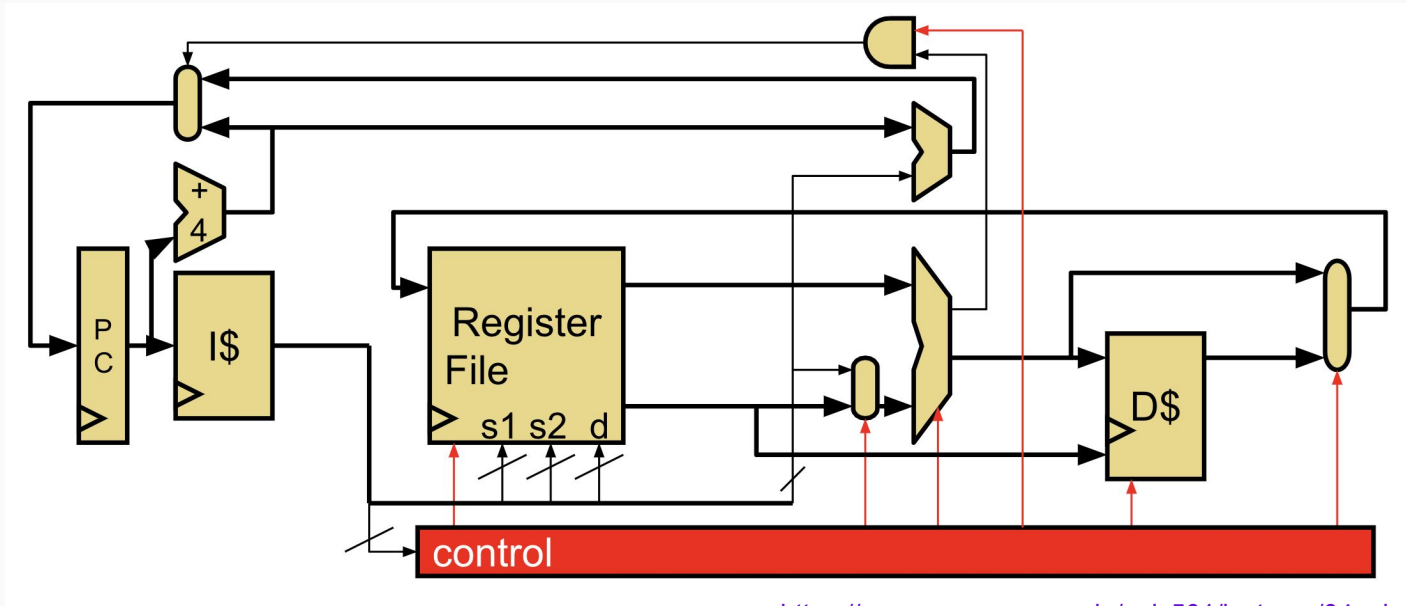
Next time: Parallel computer architecture background

CPUs have made phenomenal progress



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

Background - CPU architecture



https://www.seas.upenn.edu/~cis501/lectures/04_pipeline.pdf

- Von Neumann architecture
- Simple CPU stages: Fetch -> Decode -> Execute -> Memory -> Writeback

The performance equation

$$\text{Latency} = \text{seconds} / \text{program} = \\ (\text{instructions} / \text{program}) * (\text{cycles} / \text{instruction}) * (\text{seconds} / \text{cycle})$$

Instructions / program depends on program, compiler, instruction set architecture (ISA)

Cycles / instruction (CPI) depends on program, compiler, ISA, micro-architecture

Seconds / cycle (clock period) depends on micro-architecture, technology parameters

Fundamental performance quest in single core CPU:


*How to devote transistors on a chip to make a **single stream of instructions** run faster and faster*

Why latency?

- Pizza company needs to decide between delivering pizzas hot or delivering a large number of pizzas per hour
- In the desktop world, **Amdahl's law** restricts potential speed-up
- => Make the usual case faster
=> Reduce latency for a single stream

Why latency?

- Pizza company needs to decide between delivering pizzas hot or delivering a large number of pizzas per hour
- In the desktop world, **Amdahl's law** restricts potential speed-up
- => Make the usual case faster
=> Reduce latency for a single stream




Great performance ideas for CPUs

1. Memory hierarchy (caches)
2. Pipelined execution
3. Branch prediction
4. Out-of-order execution
5. Superscalar
6. Vector processing
7. Multi-threading
8. Multi-core

1. Caching

Core count doubling ~ every 2 years

DRAM DIMM capacity doubling ~ every 3 years



Source: Lim et al., ISCA 2009.


Keep more frequent data closer

Take advantage of locality: Spatial, temporal

Costs: Large on-chip area

	Latency	Bandwidth	Size
SRAM (L1, L2, L3)	1-2ns	200GBps	1-20MB
DRAM (memory)	70ns	20GBps	1-20GB
Flash/SSD (disk)	70-90μs	200-500MBps	100-1000GB
HDD (disk)	10ms	1-150MBps	500-3000GB

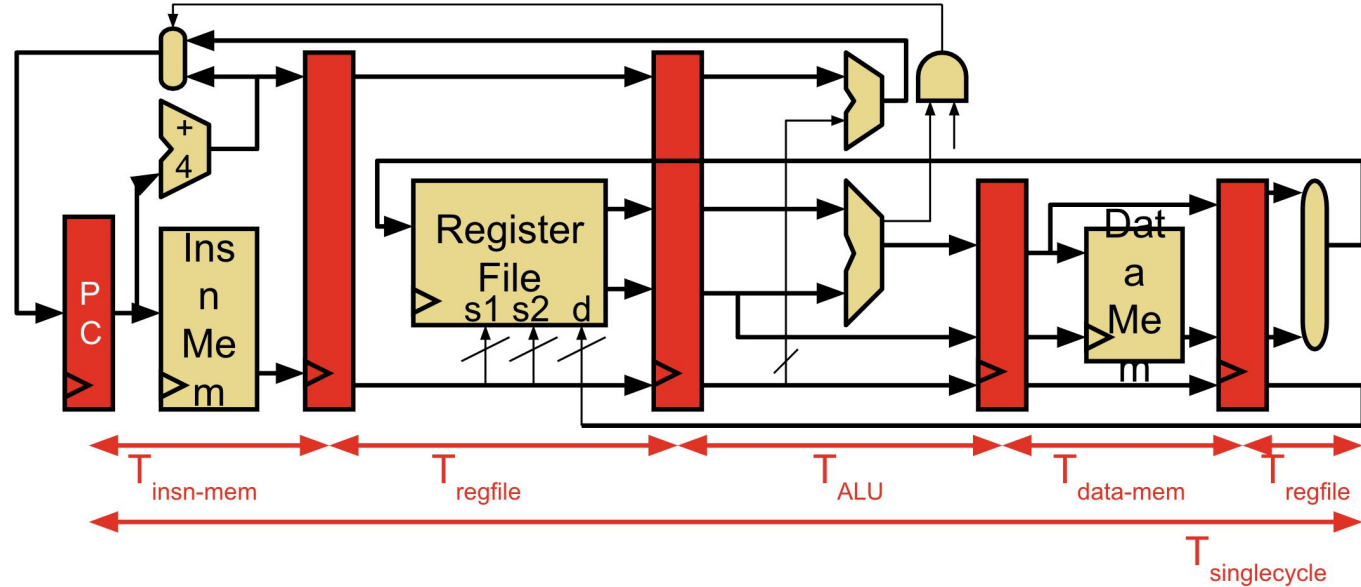
1. Caching



Intel Core i7

www.lostcircuits.com

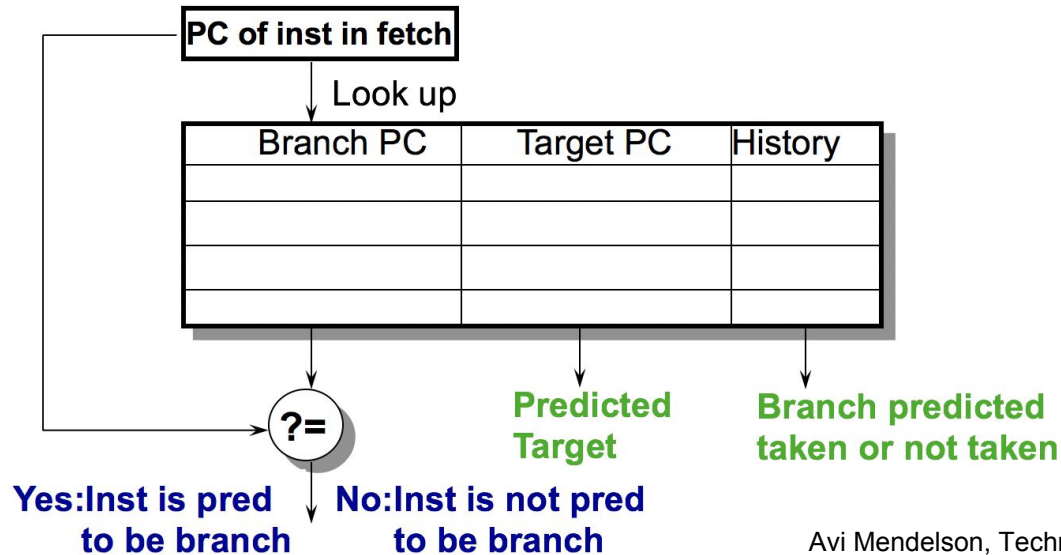
2. Pipelining



https://www.seas.upenn.edu/~cis501/lectures/04_pipeline.pdf

- Increases clocks / sec
- Costs: Larger chip area / transistor count

3. Branch prediction



Avi Mendelson, Technion

- Pipelines can stall if branch instruction arises. Branch prediction to 'guess' the branch and preemptively execute. In case of error, roll-back
- Costs: Large chip area and transistors for prediction logic, roll-back mechanism


4. Out of order execution

```
IMUL R3 ← R1, R2
ADD  R3 ← R3, R1
ADD  R4 ← R7, R6
```

```
LD   R3 ← R1(0)
ADD  R3 ← R3, R1
ADD  R4 ← R7, R6
```

- Move dependent instructions out of the way of dependent instructions
- In the above examples, instruction 2 stalls in-order pipeline
- Execute instructions **dynamically** based on when required sources are available
- Think of what can be done in compile-time and what is run-time on muArch
- Costs: Significant hardware costs of reservation stations, reorder buffer, scheduler

5. Superscalar



- Increase instructions per cycle by executing concurrently non-dependent instructions
- Requires hardware units and hardware control logic. Limited dependence on compiler
- Cost: Transistor count, hardware control-path complexity

6. Vector processing

```
for (int i = 0; i < N; i+= 4) {  
    // in parallel  
    A[i]    = B[i]    + C[i];  
    A[i+1]  = B[i+1]  + C[i+1];  
    A[i+2]  = B[i+2]  + C[i+2];  
    A[i+3]  = B[i+3]  + C[i+3];  
}
```

- Make ALUs, registers really wide enough to support operations on vectors
- Requires hardware units and compiler support. Vector extensions to x86:
 - **SSE2**: 4-wide vector operations on Intel Pentium 4, AMD Athlon 64
 - **AVX**: 8-wide vector operations on Intel Sandy Bridge, AMD Bulldozer
- Cost: Compiler support. Hardware cost

7. Multi-threading

Plot from 2002
paper from Intel
that proposed
hyperthreading





Figure 1: Single-stream performance vs. cost

- Thread composition requires maintaining architectural state: PC, registers, stack, shared globals, partitioning of ROB, other buffers
- Cost: Hardware for maintaining and switching contexts, OS management, cache conflicts, programming complexity

8. Multi-core

Inflection point in 2004, when Intel hits the power wall

Source: “No free lunch” by Dr. Dobbs in 2005

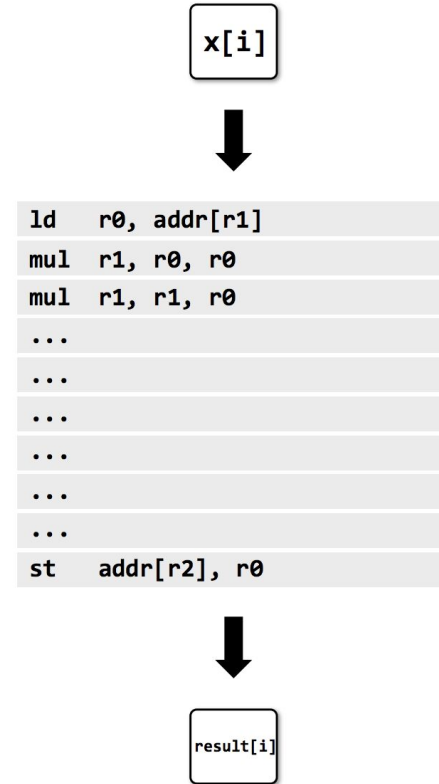



- Finally, just replicate the entire pipeline multiple times on the same chip
- No resource sharing, except for last-level (usually L3) cache
- Cost: Programmer difficulty, resource contentions

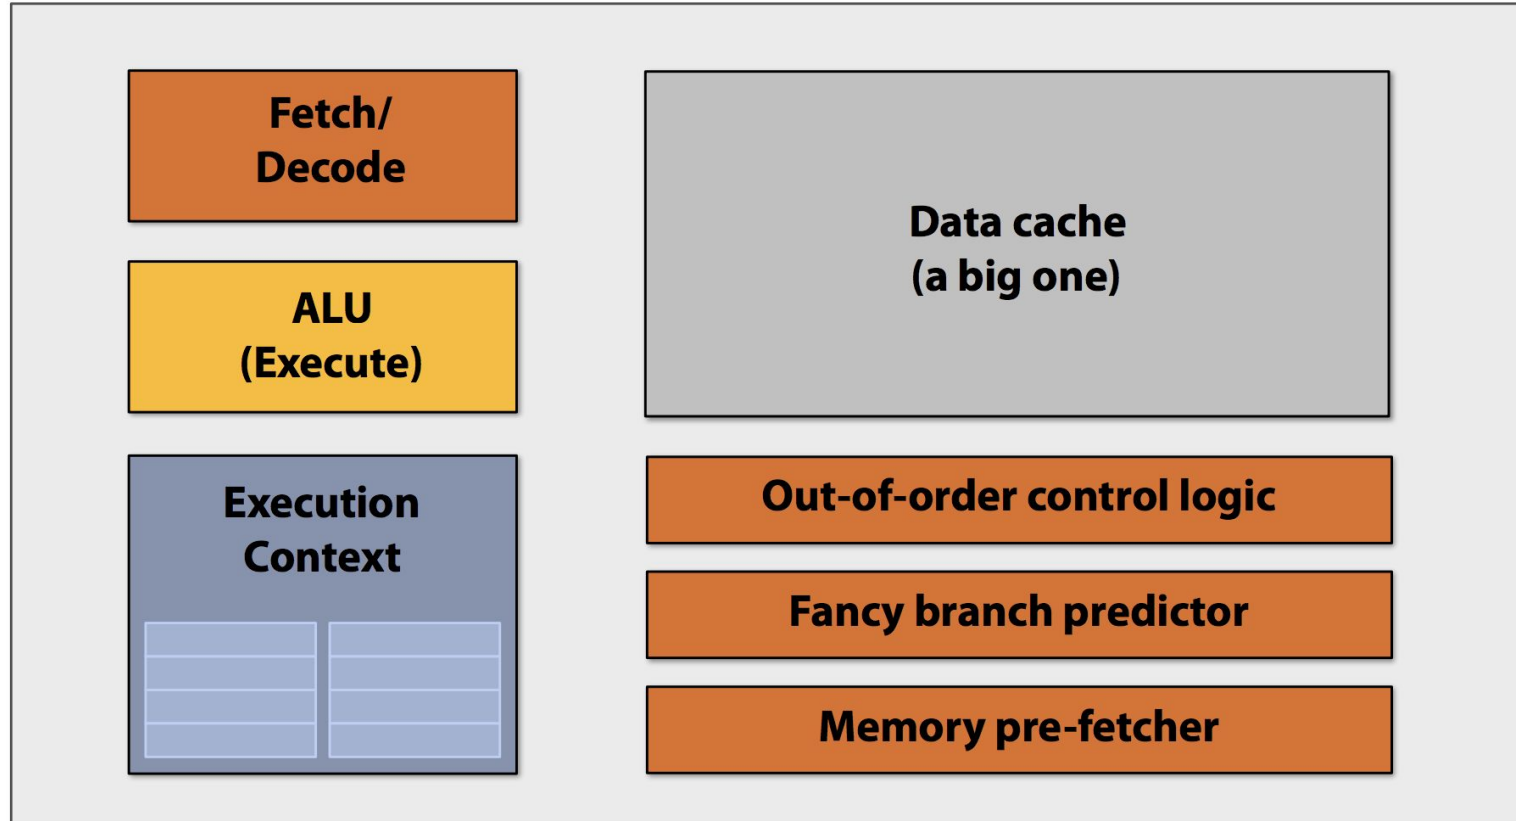
Piecing it together

- 1970s: Single stage combinational processing
- 1980s: Pipelining to exploit temporal parallelism
- 1990s: Exploit instruction-level and data-level parallelism
Out-of-order execution, superscalar, speculative execution, VLIW
- 2000s: Thread level parallelism with simultaneous multithreading (SMT)
- Mid to Late 2000s: Multi-core processing with duplicated pipelines
- 2010s: Many core processing

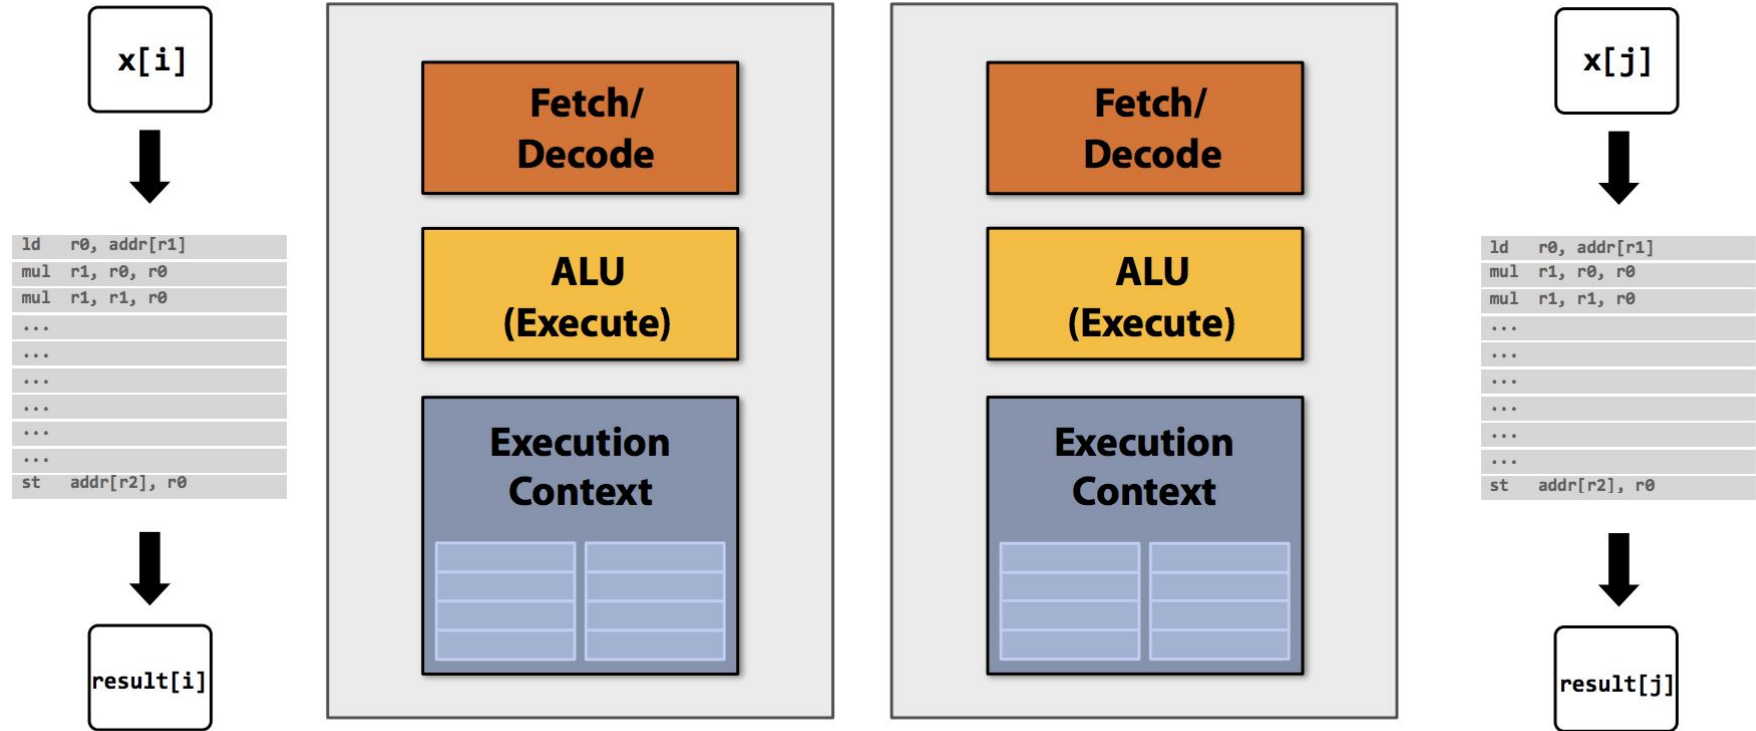
CPU - single core cartoon



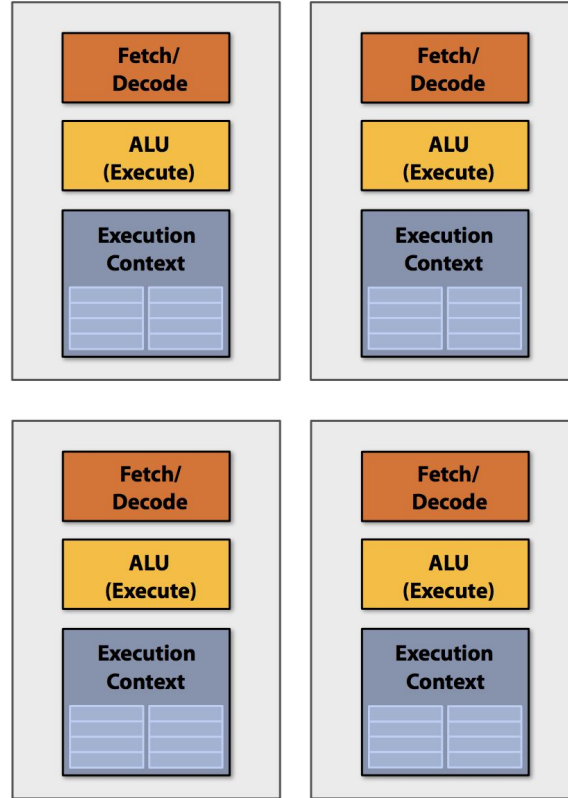

CPU cartoon - single core era



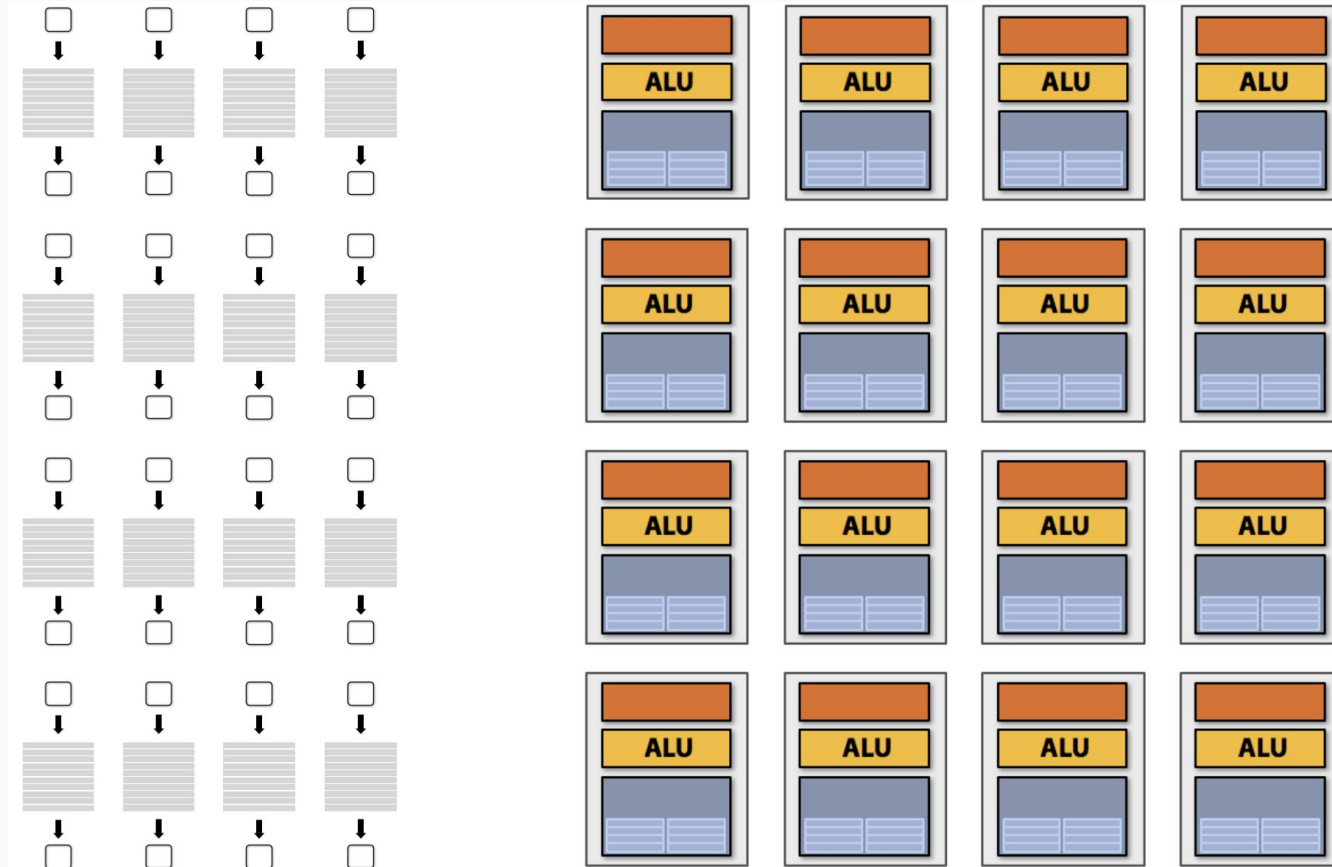
CPU cartoon - multi-core



CPU cartoon - many core



CPU cartoon - many core



Why don't we
keep doing
this for better
and better
CPUs?

Typical desktop workload

	<code>vim</code>	<code>ls</code>
Conditional branches		
Memory accesses		
Vector instructions		

Typical desktop workload

- Desktop Applications
 - Lightly threaded
 - Lots of branches
 - Lots of memory accesses

	<code>vim</code>	<code>ls</code>
Conditional branches	13.6%	12.5%
Memory accesses	45.7%	45.7%
Vector instructions	1.1%	0.2%

Next time

- Parallel computer architecture - background