

CS 6023 - GPU Programming

Full CUDA Program and Hardware Mapping

08/02/2019

Agenda

- Thread hierarchies in CUDA program
- CUDA memories
- A full CUDA program
- Mapping CUDA threads to hardware

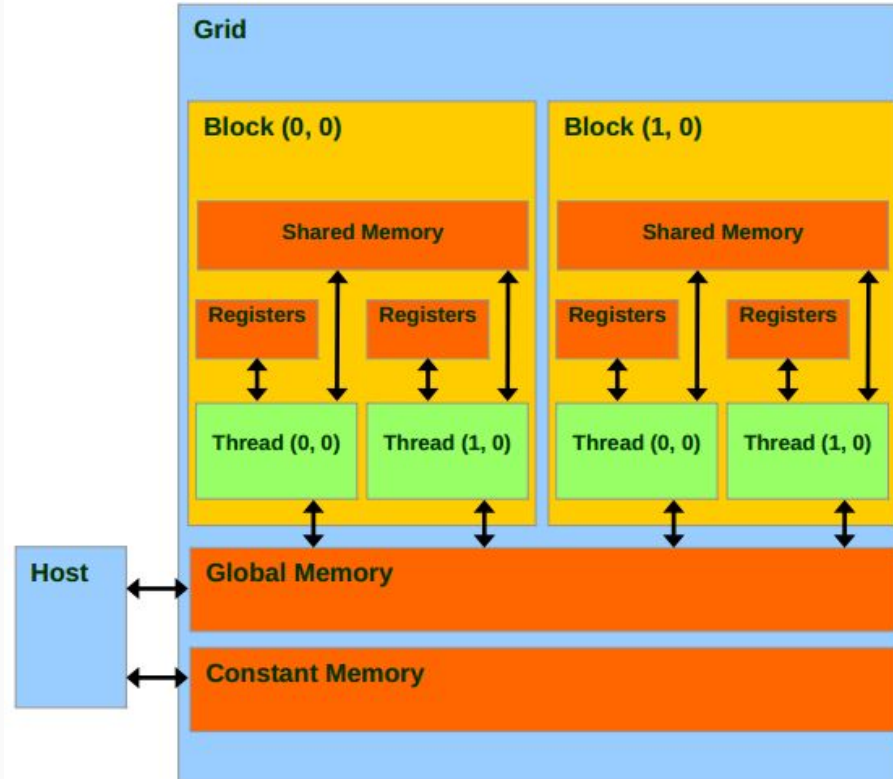
Recall: Grid/block/warp/thread

- Thread is a
- Warp is a
- Block is a
- All blocks in a grid have
- All threads in a grid share


Recall: Grid/block/warp/thread

- Thread is a sequence of instructions to be executed in SIMD manner
- Warp is a collection of threads which
 - Belong to the same block
 - Execute in SIMD fashion
- Block is a collection of threads which
 - Reside on the same processor core
 - Can share memory and synchronize
- All blocks in a grid have the same number of threads
- All threads in a grid share the same kernel function

GPU memory architecture




Global memory is huge



Green box is GK110, red lines are global memory

CUDA memory operations

Device	Thread	R/W registers
	Thread	R/W local memory
	Block	R/W shared memory
	Grid	R/W global memory
	Grid	Read only constant memory
Host	Grid	R/W global memory
	Grid	R/W constant memory



Memory allocation and freeing

```
float *deviceData = NULL;
int size = width * height * sizeof(float);

// Allocate global memory on device
cudaMalloc((void**)&deviceData, size);

// Do work

// Free memory
cudaFree(deviceData);
```


CUDA memory copy operations

```
// host to device
```

```
cudaMemcpy(devicePtr, hostPtr, size, cudaMemcpyHostToDevice);
```

```
//device to device
```

```
cudaMemcpy(destPtr, sourcePtr, size, cudaMemcpyDeviceToDevice);
```

```
//device to host
```

```
cudaMemcpy(hostPtr, devicePtr, size, cudaMemcpyDeviceToHost);
```

Compiling programs

```
__global__ void mykernel(void) {  
}  
  
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!\n");  
    return 0;  
}
```

hello_world.cu

```
$ nvcc hello_world.cu  
$ ./a.out  
Hello World!  
$
```

CUDA program for vector addition

```
int main() {  
  
    // 1. Allocate and initialize vectors A, B, and result vector C on host  
    // Allocate memory and transfer A, B, C to device  
  
    // 2. Perform  $C = A + B$  on device  
  
    // 3. Copy C from device to host  
    // Free up memory on host and device  
  
    return 0;  
}
```

CUDA program for vector addition

```
// Device code
__global__ void VecAdd(float* A, float* B, float*
C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```

```
// Host code
int main() {
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host
memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...

    // Allocate vectors in device memory
    float* d_A;    cudaMalloc(&d_A, size);
    float* d_B;    cudaMalloc(&d_B, size);
    float* d_C;    cudaMalloc(&d_C, size);
```

```
    // Copy vectors from host memory to device memory
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

    // Invoke kernel
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) /
threadsPerBlock;
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,
d_B, d_C, N);

    // Copy result from device memory to host memory
    // h_C contains the result in host memory
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);

    // Free host memory
    ...
}
```

CUDA program for vector addition

```
// Device code
__global__ void VecAdd(float* A, float* B, float*
C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```

```
// Host code
int main() {
    int N = ...;
    size_t size = N * sizeof(float);
```

```
    // Allocate input vectors h_A and h_B in host
    memory
```

```
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);
```

```
    // Initialize input vectors
```

```
    ...
```

```
    // Allocate vectors in device memory
```

```
    float* d_A;    cudaMalloc(&d_A, size);
    float* d_B;    cudaMalloc(&d_B, size);
    float* d_C;    cudaMalloc(&d_C, size);
```

```
    // Copy vectors from host memory to device memory
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
```

```
    // Invoke kernel
```

```
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) /
threadsPerBlock;
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,
d_B, d_C, N);
```

```
    // Copy result from device memory to host memory
    // h_C contains the result in host memory
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
```

```
    // Free device memory
```

```
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);
```

```
    // Free host memory
```

```
    ...
```

```
}
```

CUDA program for vector addition

```
// Device code
__global__ void VecAdd(float* A, float* B, float*
C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```

```
// Host code
int main() {
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host
    memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...
```

```
// Allocate vectors in device memory
float* d_A;    cudaMalloc(&d_A, size);
float* d_B;    cudaMalloc(&d_B, size);
float* d_C;    cudaMalloc(&d_C, size);
```

```
// Copy vectors from host memory to device memory
cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

// Invoke kernel
int threadsPerBlock = 256;
int blocksPerGrid = (N + threadsPerBlock - 1) /
threadsPerBlock;
VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,
d_B, d_C, N);

// Copy result from device memory to host memory
// h_C contains the result in host memory
cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

// Free device memory
cudaFree(d_A);
cudaFree(d_B);
cudaFree(d_C);

// Free host memory
...
}
```

CUDA program for vector addition

```
// Device code
__global__ void VecAdd(float* A, float* B, float*
C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}

// Host code
int main() {
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host
    memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...

    // Allocate vectors in device memory
    float* d_A;    cudaMalloc(&d_A, size);
    float* d_B;    cudaMalloc(&d_B, size);
    float* d_C;    cudaMalloc(&d_C, size);
```

```
    // Copy vectors from host memory to device memory
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

    // Invoke kernel
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) /
threadsPerBlock;
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,
d_B, d_C, N);

    // Copy result from device memory to host memory
    // h_C contains the result in host memory
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);

    // Free host memory
    ...
}
```

CUDA program for vector addition

```
// Device code
__global__ void VecAdd(float* A, float* B, float*
C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```

```
// Host code
int main() {
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host
    memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...

    // Allocate vectors in device memory
    float* d_A;  cudaMalloc(&d_A, size);
    float* d_B;  cudaMalloc(&d_B, size);
    float* d_C;  cudaMalloc(&d_C, size);
```

```
    // Copy vectors from host memory to device memory
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

    // Invoke kernel
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) /
threadsPerBlock;
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,
d_B, d_C, N);

    // Copy result from device memory to host memory
    // h_C contains the result in host memory
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);

    // Free host memory
    ...
}
```


CUDA program for vector addition

```
// Device code
__global__ void VecAdd(float* A, float* B, float*
C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```

```
// Host code
int main() {
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host
memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...

    // Allocate vectors in device memory
    float* d_A;  cudaMalloc(&d_A, size);
    float* d_B;  cudaMalloc(&d_B, size);
    float* d_C;  cudaMalloc(&d_C, size);
```

```
    // Copy vectors from host memory to device memory
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

    // Invoke kernel
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) /
threadsPerBlock;
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,
d_B, d_C, N);

    // Copy result from device memory to host memory
    // h_C contains the result in host memory
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);

    // Free host memory
    ...
}
```

CUDA program for vector addition

```
// Device code  
__global__ void VecAdd(float* A, float* B, float*  
C, int N){  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    if (i < N)  
        C[i] = A[i] + B[i];  
}
```

```
// Host code  
int main() {  
    int N = ...;  
    size_t size = N * sizeof(float);  
  
    // Allocate input vectors h_A and h_B in host  
memory  
    float* h_A = (float*)malloc(size);  
    float* h_B = (float*)malloc(size);  
  
    // Initialize input vectors  
    ...  
  
    // Allocate vectors in device memory  
    float* d_A;    cudaMalloc(&d_A, size);  
    float* d_B;    cudaMalloc(&d_B, size);  
    float* d_C;    cudaMalloc(&d_C, size);
```

```
    // Copy vectors from host memory to device memory  
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);  
  
    // Invoke kernel  
    int threadsPerBlock = 256;  
    int blocksPerGrid = (N + threadsPerBlock - 1) /  
threadsPerBlock;  
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,  
d_B, d_C, N);  
  
    // Copy result from device memory to host memory  
    // h_C contains the result in host memory  
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);  
  
    // Free device memory  
    cudaFree(d_A);  
    cudaFree(d_B);  
    cudaFree(d_C);  
  
    // Free host memory  
    ...  
}
```

CUDA program for vector addition

```
// Device code
__global__ void VecAdd(float* A, float* B, float*
C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}

// Host code
int main() {
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host
memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...

    // Allocate vectors in device memory
    float* d_A;    cudaMalloc(&d_A, size);
    float* d_B;    cudaMalloc(&d_B, size);
    float* d_C;    cudaMalloc(&d_C, size);
```

```
    // Copy vectors from host memory to device memory
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

    // Invoke kernel
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) /
threadsPerBlock;
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,
d_B, d_C, N);

    // Copy result from device memory to host memory
    // h_C contains the result in host memory
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);

    // Free host memory
    ...
}
```

CUDA program for vector addition

```
// Device code
__global__ void VecAdd(float* A, float* B, float*
C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}

// Host code
int main() {
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host
    memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...

    // Allocate vectors in device memory
    float* d_A;  cudaMalloc(&d_A, size);
    float* d_B;  cudaMalloc(&d_B, size);
    float* d_C;  cudaMalloc(&d_C, size);
```

```
    // Copy vectors from host memory to device memory
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

    // Invoke kernel
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) /
threadsPerBlock;
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A,
d_B, d_C, N);

    // Copy result from device memory to host memory
    // h_C contains the result in host memory
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);

    // Free host memory
    ...
}
```

Sequence of actions

- Allocate memory on host, initialize data
- Allocate memory on device `cudaMalloc`
- Copy memory from host to device
`cudaMemcpy(, , cudaMemcpyHostToDevice)`
- Invoke kernel to process on data on device
`Kernel<<< >>>(args)`
- Copy output data from device to host
`cudaMemcpy(, , cudaMemcpyDeviceToHost)`

Error handling

```
cudaError_t err = cudaMalloc((void **) &d_A, size);
if (err != cudaSuccess) {
    printf("%s in %s at line %d\n", cudaGetErrorString(err), __FILE__, __LINE__);
    exit(EXIT_FAILURE);
}
```

```
#define myCudaCheck() {
    cudaError_t err = cudaGetLastError();
    if(err != cudaSuccess) {
        printf("%s in %s at line %d\n",
            cudaGetErrorString(err), __FILE__, __LINE__);
        exit(EXIT_FAILURE);
    }
}
```



Function types

	Executed on the:	Only callable from the:
<code>__global__ void KernelFunc()</code>	device	host
<code>__device__ float DeviceFunc()</code>	device	device
<code>__host__ float HostFunc()</code>	host	host

- `__global__` must return void
- `__device__` is inlined by default
- `__host__` and `__device__` can be used together

Blocks vs Threads

Only Threads



Blocks vs Threads

Only Threads

```
// Invoke kernel  
VecAdd<<<1, N>>>(d_A, d_B, d_C, N);
```

```
// Device code  
__global__ void VecAdd(float* A, float* B, float* C, int N){  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    if (i < N)  
        C[i] = A[i] + B[i];  
}
```

Blocks vs Threads

Only Blocks

```
// Invoke kernel
```

```
// Device code
```

Blocks vs Threads



Only Blocks

```
// Invoke kernel  
VecAdd<<<N, 1>>>(d_A, d_B, d_C, N);
```

```
// Device code  
__global__ void VecAdd(float* A, float* B, float* C, int N){  
    int i = blockIdx.x;  
    if (i < N)  
        C[i] = A[i] + B[i];  
}
```

Blocks vs Threads

Blocks + Threads



Blocks vs Threads



Blocks + Threads

```
// Invoke kernel  
VecAdd<<<(N+256-1)/256, 256>>>(d_A, d_B, d_C, N);
```

```
// Device code  
__global__ void VecAdd(float* A, float* B, float* C, int N){  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    if (i < N)  
        C[i] = A[i] + B[i];  
}
```

Blocks vs Threads

Blocks + Threads (Variable threads per block)



Blocks vs Threads

Blocks + Threads (Variable threads per block)

```
// Invoke kernel
int threadsPerBlock = 256;
int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;
VecAdd<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, N);
```

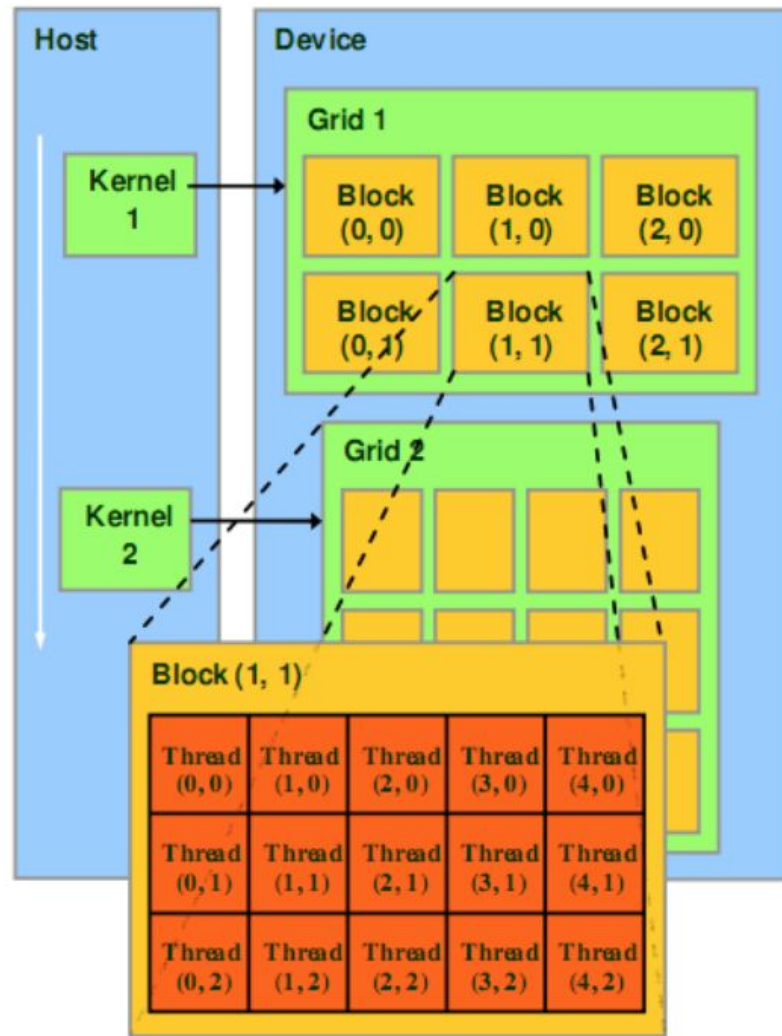
```
// Device code
__global__ void VecAdd(float* A, float* B, float* C, int N){
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```

Recap on thread hierarchies

Kernel function corresponds to a grid
which is a collection of blocks and
each block is a collection of threads

Why these multiple hierarchies?


*To map to typical high-dimensional
data structures and multiple loops in
algorithms*



Allocation of threads to SMs

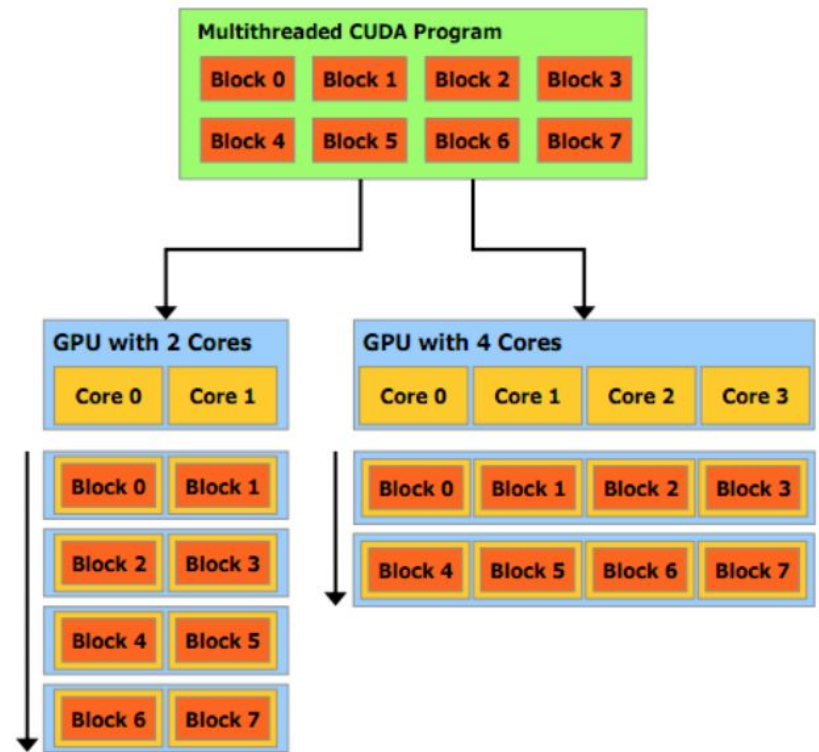
- Threads of a block execute on the same streaming multiprocessor (SM)
- Multiple blocks could be executed on the same SM
- Up to 8 blocks to each SM as resource allows
 - Fermi SM can take up to 1536 threads
 - Could be $256 \text{ (threads/block)} * 6 \text{ blocks}$
 - Or $512 \text{ (threads/block)} * 3 \text{ blocks}$, etc.

Map earlier cores to SMs
Map earlier ALUs to SP



Transparent or automatic scalability

- SM maintains thread/block idx #s
- SM manages/schedules thread execution
- CUDA abstracts the allocation information from the programmer to enable transparent or automatic scalability



A multithreaded program is partitioned into blocks of threads that execute independently from each other, so that a GPU with more cores will automatically execute the program in less time than a GPU with fewer cores.

Figure 1-4. Automatic Scalability

Recollect: Warp scheduling

- All threads in a warp execute the same instruction (recollect control divergence)
- To hide latency in scheduling warps, need almost zero overhead context switching between warps in SM
- Dynamic scheduling of warps: Warps whose next instruction has its operands ready for consumption are eligible for execution
- Eligible Warps are selected for execution based on a prioritized scheduling policy not exposed to the CUDA programmer


Warps in Hardware

- Each block is executed as **warps** of 32 threads
- This (32) is an implementation choice in the current GPU architectures, and is not available for configuration to a CUDA programmer
- Future GPUs may have different number of threads in each warp

Warps example

Consider 3 blocks assigned to an SM, with each block having 256 threads

How many warps are there in an SM?



Warps example

Consider 3 blocks assigned to an SM, with each block having 256 threads

How many warps are there in an SM?


Each block has 256

threads = $256/32 = 8$

warps

Each SM has 3 blocks =

$3 * 8 = 24$ warps



How to choose block and thread granularity

- Consider a large matrix-matrix addition operation. How to choose block size?

Recollect `VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, N);`

- For Fermi architecture: 1536 threads per SM (48 warps of 32 threads each) and up to 8 blocks per SM
- 8x8 blocks \Rightarrow
- 16x16 blocks \Rightarrow
- 32x32 blocks \Rightarrow

How to choose block and thread granularity

- Consider a large matrix-matrix addition operation. How to choose block size?

Recollect `VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, N);`

- For Fermi architecture: 1536 threads per SM (48 warps of 32 threads each) and up to 8 blocks per SM
- 8x8 blocks $\Rightarrow \leq 64$ threads per block $\Rightarrow \leq 512$ threads per SM \Rightarrow Under-utilization
- 16x16 blocks \Rightarrow
- 32x32 blocks \Rightarrow

How to choose block and thread granularity

- Consider a large matrix-matrix addition operation. How to choose block size?

Recollect `VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, N);`

- For Fermi architecture: 1536 threads per SM (48 warps of 32 threads each) and up to 8 blocks per SM
- 8x8 blocks $\Rightarrow \leq 64$ threads per block $\Rightarrow \leq 512$ threads per SM \Rightarrow Under-utilization
- 16x16 blocks $\Rightarrow \leq 256$ threads per block $\Rightarrow \leq 6$ blocks, ≤ 1536 threads per SM \Rightarrow Full utilization
- 32x32 blocks \Rightarrow

How to choose block and thread granularity

- Consider a large matrix-matrix addition operation. How to choose block size?
Recollect `VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, N);`
- For Fermi architecture: 1536 threads per SM (48 warps of 32 threads each) and up to 8 blocks per SM
- 8x8 blocks $\Rightarrow \leq 64$ threads per block $\Rightarrow \leq 512$ threads per SM \Rightarrow Under-utilization
- 16x16 blocks $\Rightarrow \leq 256$ threads per block $\Rightarrow \leq 6$ blocks, ≤ 1536 threads per SM \Rightarrow Full utilization
- 32x32 blocks $\Rightarrow \leq 1024$ threads per block $\Rightarrow 1$ block per SM \Rightarrow Under-utilization

Next time

- Memory and related optimization of CUDA programs